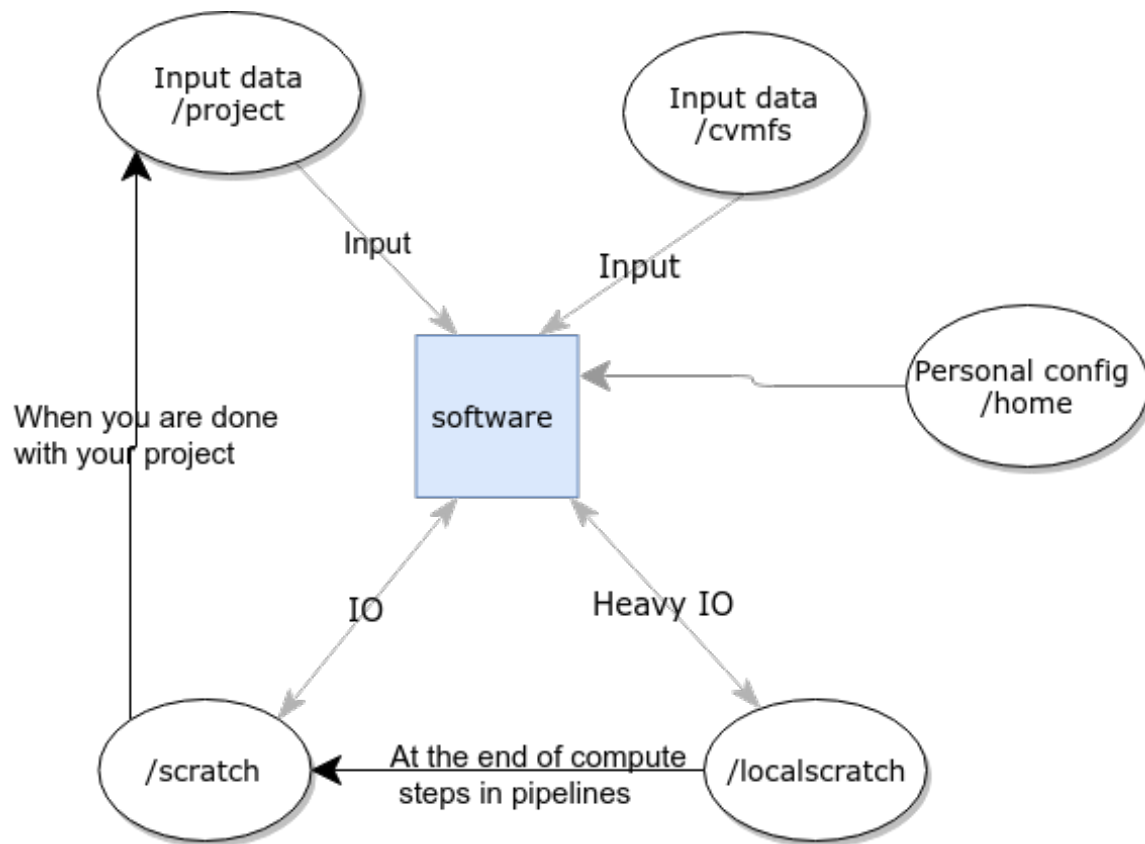


Good HPC Data Flow



Guidelines:

- Inputs are loaded from `/project` and `/cvmfs`.
- Your personal configs are stored somewhere in your `$HOME` folder.
- Store your outputs in the `/scratch` space.
- If a step of the pipeline has heavy IO use `$SLURM_TMPDIR`
- Once your project is completed (and the paper is published!) clean your `/scratch` and move what you want to keep to `/project`
- Put dormant projects in a tar ball, even if you have no intent of archiving them on tape.

Of course these are general guidelines, and you should not be too rigid about them. But in the long run, they will reduce the chance that your pipelines freeze the whole system. They will also prevent you from receiving emails telling you to "clean up your stuff!" in the `/project` space.